

# The Apriori Algorithm

Association rule learning,  
the Apriori algorithm and  
it's implementation

tommyod @ github

Presentation: [github.com/tommyod/Efficient-Apriori/blob/master/docs/presentation/apriori.pdf](https://github.com/tommyod/Efficient-Apriori/blob/master/docs/presentation/apriori.pdf)

December 28, 2018

# Table of contents

A problem: learning association rules

A solution: the Apriori algorithm

A practical matter: writing a Python implementation

Summary and references

A problem: learning association rules

# Motivating example

## Example (Learning from transactions)

Consider the following set of *transactions*.

{eggs, bread, jam, bacon}

{apples, eggs, bacon}

{bacon, bread}

{ice cream, bread, bacon}

What interesting information can we infer from this data?

Examples:

- The itemsets {bacon, bread} and {bacon, eggs} often appear in the transactions, with counts 3 and 2, respectively.
- The rule {bread}  $\Rightarrow$  {bacon} is meaningful in the sense that  $P(\text{bacon}|\text{bread}) = 1$ .

# Formal problem statement

## Problem

*Given a database  $T = \{t_1, t_2, \dots, t_m\}$ , where the  $t_i$  are transactions, and a set of items  $I = \{i_1, i_2, \dots, i_n\}$ , learn meaningful rules  $X \Rightarrow Y$ , where  $X, Y \subset I$ .*

To accomplish this, we need measures of the *meaningfulness* of association rules.

# Properties of association rules

## Definition (Support)

The *support* of an association rule  $X \Rightarrow Y$  is the frequency of which  $X \cup Y$  appears in the transactions  $T$ , i.e.  $\text{support}(X \Rightarrow Y) := P(X, Y)$ .

- No reason to distinguish between the support of an itemset, and the support of an association rule, i.e.  $\text{support}(X \Rightarrow Y) = \text{support}(X \cup Y)$ .
- An important property of support is that  $\text{support}(\{\text{eggs}, \text{bacon}\}) \leq \text{support}(\{\text{bacon}\})$ .

More formally, we observe that:

## Theorem (Downward closure property of sets)

If  $s \subset S$ , then  $\text{support}(s) \geq \text{support}(S)$ .

# Properties of association rules

## Definition (Confidence)

The confidence of the association rule  $X \Rightarrow Y$  is given by

$$\text{confidence}(X \Rightarrow Y) = P(Y|X) = \frac{P(X, Y)}{P(X)} = \frac{\text{support}(X \Rightarrow Y)}{\text{support}(X)}.$$

Notice the following interesting property.

## Example

The confidence of  $\{A, B\} \Rightarrow \{C\}$  will always be greater than, or equal to,  $\{A\} \Rightarrow \{B, C\}$ ). By definition we have

$$\frac{\text{support}(\{A, B\} \Rightarrow \{C\})}{\text{support}(\{A, B\})} \geq \frac{\text{support}(\{A\} \Rightarrow \{B, C\})}{\text{support}(\{A\})},$$

where the numerator is equal, and  $\text{support}(\{A\}) > \text{support}(\{A, B\})$

# Properties of association rules

## Definition (Confidence)

The confidence of the association rule  $X \Rightarrow Y$  is given by

$$\text{confidence}(X \Rightarrow Y) = P(Y|X) = \frac{P(X, Y)}{P(X)} = \frac{\text{support}(X \Rightarrow Y)}{\text{support}(X)}.$$

## Theorem (Downward closure property of rules)

Consider the rule  $(X - y) \Rightarrow y$  and  $(X - Y) \Rightarrow Y$ , where  $y \subset Y$ . Then

$$\text{confidence}((X - y) \Rightarrow y) \geq \text{confidence}((X - Y) \Rightarrow Y)$$

**Proof.** The numerator is identical, but the denominator has  $\text{support}(X - y) \leq \text{support}(X - Y)$  by the downward closure property of sets.



# Examples of support and confidence

## Example (Support and confidence of a rule)

Consider again the following set of transactions.

{eggs, bread, jam, bacon}

{apples, eggs, bacon}

{bacon, bread}

{ice cream, bread, bacon}

- The rule  $\{\text{bread}\} \Rightarrow \{\text{bacon}\}$  has support  $3/4$ , confidence 1.
  - Support  $3/4$  since  $\{\text{bread, bacon}\}$  appears in 3 of the transactions.
  - Confidence 1 since  $\{\text{bread}\}$  appears 3 times, and in 3 of those  $\{\text{bacon}\}$  also appears.

# A naive algorithm

## Example (Naive algorithm for learning rules)

```
for subsets of every size  $k = 1, \dots, |I|$ 
  for every subset of size  $k$ 
    for every split of this subset into  $\{X\} \Rightarrow \{Y\}$ 
      compute support and confidence of the rule
      by counting the support in the transactions
```

- Fantastic starting point for an algorithm, since it (1) clearly terminates in finite time, (2) is simple to implement and (3) will run reasonably fast on small problem instances.
- Terribly slow on realistic problem instances, since it must check every possible itemset against every transaction.

A solution: the Apriori algorithm

# Overview of apriori

- Split the problem into two distinct phases.
  - Finding meaningful (high support) itemsets.
  - Generating meaningful (high confidence) rules.
- **Phase 1**
  - The user specifies a desired *minimum support*.
  - The algorithm exploits the downward closure property, i.e.  $\text{support}(S) \leq \text{support}(s)$  if  $s \subset S$ .
    - \* No reason to check  $S$  if  $s$  has low support.
  - Bottom-up approach to subset generation.
- **Phase 2**
  - The user specifies a desired *minimum confidence*.
  - Also exploits the above downward closure property.
  - Bottom-up approach to rule generation.

# Phase 1: Generating itemsets (example 1)

## Example (Itemset generation via Apriori)

Consider again the following set of transactions.

{eggs, bread, jam, bacon}

{apples, eggs, bacon}

{bacon, bread}

{ice cream, bread, bacon}

- We set the minimum confidence to 50 %.
  - Itemsets of size 1 with desired confidence are {bacon}, {bread} and {eggs}. They are called *large itemsets* of size 1.
  - From these, we can form {bacon, bread}, {bacon, eggs} and {bread, eggs}. These are *candidate itemsets* of size 2.
  - Large itemsets of size 2: {bacon, bread} and {bacon, eggs}.

## Phase 1: Generating itemsets (example 2)

### Example

#### Transactions

$\{1, 2, 7, 4\}$

$\{2, 3, 4\}$

$\{1, 6, 3\}$

$\{1, 2, 4, 5\}$

#### Iteration 1

- Running the algorithm with minimum support 50 %.
- Candidate itemsets of size 1:
  - $\{1\}, \{2\}, \{3\}, \{4\}, \{5\}, \{6\}, \{7\}$
- Large itemsets of size 1:
  - $\{1\}, \{2\}, \{3\}, \{4\}$

## Phase 1: Generating itemsets (example 2)

### Example

#### Transactions

{1, 2, 7, 4}

{2, 3, 4}

{1, 6, 3}

{1, 2, 4, 5}

#### Iteration 2

- Running the algorithm with minimum support 50 %.
- Candidate itemsets of size 2:
  - {1, 2}, {1, 3}, {1, 4}, {2, 3}, {2, 4}, {3, 4}
- Large itemsets of size 2:
  - {1, 2}, {1, 4}, {2, 4}

## Phase 1: Generating itemsets (example 2)

### Example

#### Transactions

$\{1, 2, 7, 4\}$

$\{2, 3, 4\}$

$\{1, 6, 3\}$

$\{1, 2, 4, 5\}$

#### Iteration 3

- Running the algorithm with minimum support 50 %.
- Candidate itemsets of size 3:
  - $\{1, 2, 4\}$
- Large itemsets of size 3:
  - $\{1, 2, 4\}$



# Phase 1: Pseudocode

## Algorithm sketch

Create  $L_1$ , a set of large itemsets of size 1

$j = 1$

while  $L_j$  is not empty do:

    create every candidate set  $C_{j+1}$  from  $L_j$

    prune candidates a priori  $C_{j+1}$  (every subset must be in  $L_j$ )

    for every transaction  $t_i \in T$  do:

        count occurrences of every set in  $C_{j+1}$  in  $t_i$

$j = j + 1$

---

Iterating through the transactions checking for every possible candidate in  $C_{j+1}$  is expensive. Optimizations: choosing good data structures, pruning transactions.

## Phase 1: Pseudocode - Details on candidates and pruning

create every candidate set  $C_{j+1}$  from  $L_j$

prune candidates a priori  $C_{j+1}$  (every subset must be in  $L_j$ )

---

**Example** Given large itemsets of size 3

$\{1, 2, 3\}, \{1, 2, 4\}, \{1, 3, 4\}, \{1, 3, 5\}, \{2, 3, 4\}$ .

- Naive candidates are  $\{2, 3, 4, 5\}, \{1, 3, 4, 5\}, \{1, 2, 4, 5\}, \{1, 2, 3, 5\}, \{1, 2, 3, 4\}$ .
- Apriori-gen candidates are  $\{1, 2, 3, 4\}, \{1, 3, 4, 5\}$ . Generated efficiently by keeping the itemsets sorted.
- While the itemset  $\{1, 2, 3, 4\}$  is kept,  $\{1, 3, 4, 5\}$  is discarded since the subset  $\{1, 3, 5\} \subset \{1, 3, 4, 5\}$  is not among the large itemsets of size 3 .

The example above is from page 4 in the referenced paper.

## Phase 1: Pseudocode - Details on counting occurrences

for every transaction  $t_i \in T$  do:  
    count occurrences of every set in  $C_{j+1}$  in  $t_i$

---

### Example

Check if  $A = \{1, 3, 7\}$  is a subset of  $B = \{1, 2, 3, 5, 7, 9\}$ .

- A naive computation checks if every element of  $A$  is found in  $B$ . This has computational complexity  $\mathcal{O}(|A||B|)$ , where  $|A|$  is the size of  $A$ .
- A better approach is to use binary search when  $B$  is sorted. The computational complexity becomes  $\mathcal{O}(|A| \log_2 |B|)$ .
- Using hash tables (e.g. the built-in `set.issubset` in Python), the computational complexity is down to  $\mathcal{O}(|A|)$ .

For the given example, this resolves to approximately 18, 8 and 3 operations.

## Phase 2: Building association rules (example)

- In practice this step is much faster than Phase 1.
- The efficient algorithm exploits the downward closure property.

### Example

Consider rules made from  $ABCD$ . First the algorithm tries to move itemsets of size 1 to the right hand side, i.e. one of  $\{\{A\}, \{B\}, \{C\}, \{D\}\}$ .

$$\begin{array}{ll} BCD \Rightarrow A & ACD \Rightarrow B \\ ABD \Rightarrow C & ABC \Rightarrow D \end{array}$$

Assume that only  $ABC \Rightarrow D$  and  $ACD \Rightarrow B$  had high enough confidence. Then the only rule created from  $ABCD$  with a size 2 itemset on the right hand side worth considering is  $AC \Rightarrow BD$ . This is a direct result of the downward closure property.

Recursive function which is not very easy to explain in detail.

# The Apriori algorithm on real data

Consider the following data set, with 32.561 rows.

Education	Marital-status	Relationship	Race	Sex	Income	Age
Bachelors	Never-married	Not-in-family	White	Male	$\leq 50K$	middle-aged
Bachelors	Married-civ-spouse	Husband	White	Male	$\leq 50K$	old
HS-grad	Divorced	Not-in-family	White	Male	$\leq 50K$	middle-aged
11th	Married-civ-spouse	Husband	Black	Male	$\leq 50K$	old
Bachelors	Married-civ-spouse	Wife	Black	Female	$\leq 50K$	young
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
Masters	Married-civ-spouse	Wife	White	Female	$\leq 50K$	middle-aged
9th	Married-spouse-absent	Not-in-family	Black	Female	$\leq 50K$	middle-aged
HS-grad	Married-civ-spouse	Husband	White	Male	$> 50K$	old
Masters	Never-married	Not-in-family	White	Female	$> 50K$	middle-aged

The data may be found at <https://archive.ics.uci.edu/ml/datasets/adult>.

# The Apriori algorithm on real data

Some rules are obvious in retrospect:

$$\{\text{Husband}\} \Rightarrow \{\text{Male}\}$$

$$\{\leq 50\text{K}, \text{Husband}\} \Rightarrow \{\text{Male}\}$$

$$\{\text{Husband}, \text{middle-aged}\} \Rightarrow \{\text{Male}, \text{Married-civ-spouse}\}$$

Some are more interesting:

$$\{\text{HS-grad}\} \Rightarrow \{\leq 50\text{K}\}$$

$$\{\leq 50\text{K}, \text{young}\} \Rightarrow \{\text{Never-married}\}$$

$$\{\text{Husband}\} \Rightarrow \{\text{Male}, \text{Married-civ-spouse}, \text{middle-aged}\}$$

The meaningfulness of a rule may be measured by *confidence*, *lift* and *conviction*.

A practical matter: writing a Python implementation

# Overview of workflow

- Write simple functions first, i.e. the building blocks (e.g. pruning)
- Add doctests and unit tests (e.g. examples from paper)
- Implement a naive, but correct algorithm
- Implement an asymptotically fast algorithm
- Test the preceding two implementations against each other
- Optimize implementation by profiling the code (find bottlenecks)

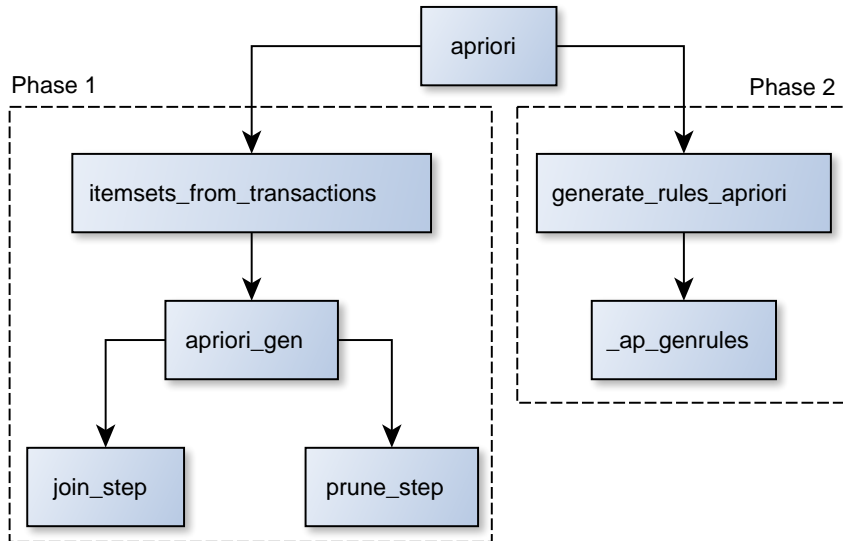
Understand → Naive algorithm → Asymptotically fast → Further optimizations



# Software testing

- Unit tests
  - Test a simple function  $f(x_i) = y_i$  for known cases  $i = 1, 2, \dots$
  - Doubles as documentation when writing *doctests* in Python
- Property tests
  - Fix a property, i.e.  $f(a, b) = f(b, a)$  for every  $a, b$
  - Generate many random inputs  $a, b$  to make sure the property holds
- Testing against R, Wikipedia, etc
  - Generate some inputs and test against the *arules* package

# Software structure



Software found at <https://github.com/tommyod/Efficient-Apriori>.

## Summary and references

## Summary and references

The Apriori algorithm discovers frequent itemsets in phase 1, and meaningful association rules in phase 2. Both phases employ clever bottom-up algorithms. By application of the downward closure property of itemsets (support) and rules (confidence), candidates may be pruned prior to expensive computations.

- The Python implementation
  - [github.com/tommyod/Efficient-Apriori](https://github.com/tommyod/Efficient-Apriori)
- The original paper
  - Agrawal et al, *Fast Algorithms for Mining Association Rules*, 1994  
<http://www.cse.msu.edu/~cse960/Papers/MiningAssoc-AgrawalAS-VLDB94.pdf>